

## DATA INSPECTION/CLEANING

```
data = pd.read_csv('/Users/tonydao/Documents/PythonProjects/LungCancerProject/lungCancer.csv')
print(data.head())
print(data.info())
```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	\
0	M	69	1	2	2	1	
1	M	74	2	1	1	1	
2	F	59	1	1	1	2	
3	M	63	2	2	2	1	
4	F	63	1	2	1	1	

	CHRONIC_DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL_CONSUMING	COUGHING	\
0	1	2	2	1	2	2	
1	2	2	2	2	1	1	
2	1	2	1	2	1	2	
3	1	1	1	1	2	1	
4	1	1	1	2	1	2	

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN	LUNG_CANCER
0	2	2	2	YES
1	2	2	2	YES
2	2	1	2	NO
3	1	2	2	NO
4	2	1	1	NO

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 309 entries, 0 to 308
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   GENDER                                309 non-null    object
1   AGE                                  309 non-null    int64
2   SMOKING                              309 non-null    int64
3   YELLOW_FINGERS                       309 non-null    int64
4   ANXIETY                              309 non-null    int64
5   PEER_PRESSURE                        309 non-null    int64
6   CHRONIC_DISEASE                      309 non-null    int64
7   FATIGUE                              309 non-null    int64
8   ALLERGY                              309 non-null    int64
9   WHEEZING                             309 non-null    int64
10  ALCOHOL_CONSUMING                    309 non-null    int64
11  COUGHING                             309 non-null    int64
12  SHORTNESS OF BREATH                  309 non-null    int64
13  SWALLOWING DIFFICULTY                309 non-null    int64
14  CHEST PAIN                           309 non-null    int64
15  LUNG_CANCER                          309 non-null    object
dtypes: int64(14), object(2)
memory usage: 38.8+ KB
None
```

- Here we can see what cols we are working with and the data type that exist within each cols.
- We have 16 cols.
  - 2 are type of Objects
  - 14 are floats.
- I would change any funky label names with something simple.

### • Checking for missing values and duplicates

```
print(data.isnull().sum())
print(data.duplicated().sum())
```

GENDER	0
AGE	0
SMOKING	0
YELLOW_FINGERS	0
ANXIETY	0
PEER_PRESSURE	0
CHRONIC_DISEASE	0
FATIGUE	0
ALLERGY	0
WHEEZING	0
ALCOHOL_CONSUMING	0
COUGHING	0
SHORTNESS OF BREATH	0
SWALLOWING DIFFICULTY	0
CHEST PAIN	0
LUNG_CANCER	0
dtype: int64	
33	

- Here we see non of our cols contain any nulls.
- We do have 33 duplicates.

```

duplicates = data[data.duplicated()]
print(duplicates)

```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE
99	M	56	2	1	1	1
100	M	58	2	1	1	1
117	F	51	2	2	2	2
199	F	55	2	1	1	2
212	M	58	2	1	1	1
223	M	63	2	2	2	1
256	M	60	2	1	1	1
275	M	64	2	2	2	2
284	M	58	2	2	2	2
285	F	58	2	2	2	2
286	F	63	1	1	1	1
287	F	51	2	2	2	2
288	F	61	1	2	2	2
289	F	61	2	1	1	1
290	M	76	2	1	1	1
291	M	71	2	2	2	1
292	M	69	1	1	2	1
293	F	56	2	2	2	1
294	M	67	1	1	1	2
295	F	54	2	2	2	1
296	M	63	1	2	1	1
297	F	47	2	2	1	2
298	M	62	2	1	2	1
299	M	65	2	2	2	2
300	F	63	2	2	2	2
301	M	64	1	2	2	2
302	F	65	2	2	2	2
303	M	51	1	2	1	1
304	F	56	1	1	1	2
305	M	70	2	1	1	1
306	M	58	2	1	1	1
307	M	67	2	1	2	1
308	M	62	1	1	1	2

- If we ever need the duplicates ones only then this variable holds it and we can display the specific duplicated rows.

```

data = data.drop_duplicates()

print(data.duplicated().sum())

0

```

- We drop any duplicated rows, and we can confirm our results.

## • Standardize Categorical Data

- Our 'LUNG\_CANCER' col is 'YES/NO'. Since all of our others rows are integer. We can represent this cols as '1/0' to allow data analysis be conducted more efficiently.

```

data['LUNG_CANCER'] = data['LUNG_CANCER'].map({'YES': 1, 'NO': 0})

```

/var/folders/g7/k8tb4tqx737cssc6482hmgb40000gn/T/ipykernel\_1593/261  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead  
  
See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min/05min.html#update-in-place>  
data['LUNG\_CANCER'] = data['LUNG\_CANCER'].map({'YES': 1, 'NO': 0})

```

data['LUNG_CANCER']

```

0	1
1	1
2	0
3	0
4	0
...	...
279	1
280	0
281	0
282	0
283	1

Name: LUNG\_CANCER, Length: 276, dtype: int64

## BONUS TIP:

- Remove any trailing whitespace.
- Convert the cols that are considered as "categorical" to type "category".
  - reduces memory usage.
  - allows more efficient data manipulation and plotting.

```
data.columns = data.columns.str.strip()
categorical_columns = [
    'GENDER', 'SMOKING', 'YELLOW_FINGERS', 'ANXIETY', 'PEER_PRESSURE',
    'CHRONIC_DISEASE', 'FATIGUE', 'ALLERGY', 'WHEEZING', 'ALCOHOL_CONSUMING',
    'COUGHING', 'SHORTNESS_OF_BREATH', 'SWALLOWING_DIFFICULTY', 'CHEST_PAIN'
]

for col in categorical_columns:
    data[col] = data[col].astype('category')

data.info()


<class 'pandas.core.frame.DataFrame'>
Index: 276 entries, 0 to 283
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   GENDER                276 non-null   category
1   AGE                   276 non-null   int64
2   SMOKING               276 non-null   category
3   YELLOW_FINGERS        276 non-null   category
4   ANXIETY               276 non-null   category
5   PEER_PRESSURE         276 non-null   category
6   CHRONIC_DISEASE       276 non-null   category
7   FATIGUE               276 non-null   category
8   ALLERGY               276 non-null   category
9   WHEEZING              276 non-null   category
10  ALCOHOL_CONSUMING     276 non-null   category
11  COUGHING              276 non-null   category
12  SHORTNESS_OF_BREATH   276 non-null   category
13  SWALLOWING_DIFFICULTY 276 non-null   category
14  CHEST_PAIN            276 non-null   category
15  LUNG_CANCER           276 non-null   int64
dtypes: category(14), int64(2)
memory usage: 11.9 KB
```

## • Quick Data Visualization


```
# Target variable
data['LUNG_CANCER'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Distribution of Lung Cancer Diagnosis')
plt.xlabel('Lung Cancer (0=No, 1=Yes)')
plt.ylabel('Count')
plt.show()

# Age distribution
plt.hist(data['AGE'], bins=20, color='lightgreen', edgecolor='black')
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()


# Gender distribution
data['GENDER'].value_counts().plot(kind='bar', color=['purple', 'orange'])
plt.title('Gender Distribution')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.show()
```



- Strong imbalance of Cancer to no-cancer



- Age group is fairly normally distributed with most patients between 50-75 years old.



- Genders are fairly balanced.

## • Advance Data Inspection

### ○ Multivariate Outlier Detection

```
import numpy as np
from sklearn.ensemble import IsolationForest
data.columns = data.columns.str.strip()
# Correct: Create a copy of the original data
outlier_data = data.copy()

# Convert categorical columns to numeric codes
for col in categorical_columns:
    data[col] = data[col].astype('category') # Convert to category FIRST
    print(f'{col} dtype: {data[col].dtype}') # Verify conversion
    outlier_data = data.copy()
    for col in categorical_columns:
        outlier_data[col] = outlier_data[col].cat.codes

GENDER dtype: category
SMOKING dtype: category
YELLOW_FINGERS dtype: category
ANXIETY dtype: category
PEER_PRESSURE dtype: category
CHRONIC_DISEASE dtype: category
FATIGUE dtype: category
ALLERGY dtype: category
WHEEZING dtype: category
ALCOHOL_CONSUMING dtype: category
COUGHING dtype: category
SHORTNESS OF BREATH dtype: category
SWALLOWING DIFFICULTY dtype: category
CHEST PAIN dtype: category
```

```
# Verify conversion
print(outlier_data[categorical_columns].head())

# Detect outliers (exclude target variable)
iso = IsolationForest(contamination=0.05, random_state=42)
outliers = iso.fit_predict(outlier_data.drop('LUNG_CANCER', axis=1))

# Add outlier flags to original data
data['OUTLIER_FLAG'] = np.where(outliers == -1, 1, 0)

# Analyze outliers
outlier_profile = data[data['OUTLIER_FLAG'] == 1].describe().T
print("Outlier Summary:\n", outlier_profile[['mean', 'std', 'min', 'max']])
```

	GENDER	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC_DISEASE
0	1	0	1	1	0	0
1	1	1	0	0	0	1
2	0	0	0	0	1	0
3	1	1	1	1	0	0
4	0	0	1	0	0	0

	FATIGUE	ALLERGY	WHEEZING	ALCOHOL_CONSUMING	COUGHING	\
0	1	0	1	1	1	
1	1	1	0	0	0	
2	1	0	1	0	1	
3	0	0	0	1	0	
4	0	0	1	0	1	

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN
0	1	1	1
1	1	1	1
2	1	0	1
3	0	1	1
4	1	0	0

- Outliers cluster around mean age 62 +/- 7.3 years
- Age Range: 47-77 years
- Age only continuous variable.
- Any outliers are flagged as 1.

### ○ Advance Numerical Analysis

	mean	std	min	max
AGE	62.642857	7.344356	47.0	77.0
LUNG_CANCER	0.642857	0.497245	0.0	1.0
OUTLIER_FLAG	1.000000	0.000000	1.0	1.0

```

from scipy.stats import kendalltau

# Create symptom score dynamically from existing columns
symptoms = ['COUGHING', 'SHORTNESS OF BREATH', 'CHEST PAIN',
            'SWALLOWING DIFFICULTY', 'WHEEZING']

# Convert to numerical codes
for col in symptoms:
    data3[col] = data3[col].astype('category') # Ensure categorical dtype
    data3[col] = data3[col].cat.codes.astype('int64') # Force int64
# Create symptom score with int64 dtype
data3['SYMPTOM_SCORE'] = data3[symptoms].sum(axis=1).astype('int64')

# Ensure target is int64
data3['LUNG_CANCER'] = data3['LUNG_CANCER'].astype('int64')

int8_cols = data3.select_dtypes(include=['int8']).columns
data3[int8_cols] = data3[int8_cols].astype('int64')

# Analyze non-parametric correlations
num_features = ['AGE', 'SYMPTOM_SCORE']
target = data['LUNG_CANCER']

corr_results = []
for feat in num_features:
    tau, p_value = kendalltau(data[feat], target)
    corr_results.append({
        'Feature': feat,
        'Kendall Tau': tau,
        'p-value': p_value
    })

corr_df = pd.DataFrame(corr_results)
print("Non-Parametric Correlation:\n", corr_df)

```

	Feature	Kendall Tau	p-value
0	AGE	0.080379	1.092943e-01
1	SYMPTOM_SCORE	0.339395	4.183125e-10

- Age vs Lung Cancer
  - Our Kendal Tau value of 0.08 shows a weak positive relationship
  - $p = 0.109$ , thus it is not statistically sig.
- Symptom Score vs Lung Cancer
  - Kendal Tau = 0.34
    - moderate positive relationship
  - $p = \text{less than } 4.18 \times 10^{-10}$ 
    - highly statistically sig.
- For every 1 unit increase, there is 34% increased likelihood of lung cancer.
- Odd Ratio =  $1 + \tau / 1 - \tau = 2.03$ 
  - So 1 unit increase = 103% increase in **ODDS**

## EXPLORATORY DATA ANALYSIS

### • Target Variable Analysis

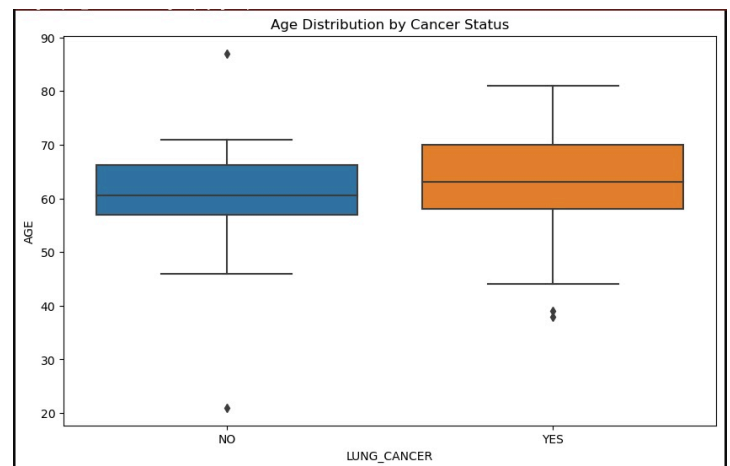
```

import seaborn as sns
# Class distribution with detailed metrics
target_dist = data3['LUNG_CANCER'].value_counts()
print(f"Class Balance:\n{target_dist}")
print(f"\nPositive Class Percentage: {target_dist[1]/len(data3)*100:.1f}%")

Class Balance:
LUNG_CANCER
1    238
0     38
Name: count, dtype: int64

Positive Class Percentage: 86.2%

```



- We have 238 cases (86.2%) that are pos for lung cancer.
- 39 cases (13.8%) are negative.
- From the boxplot
  - both groups have similar median ages. Ranges and outliers differ.
  - IQR is wider for pos cases.
  - Median Line: 50% of cancer patients are older than 62, while 50% of non-cancer patients are older than 56.



```
counts = data3['LUNG_CANCER'].value_counts()
percentages = counts / len(data3) * 100
print("Class distribution:\n", counts)
print("Class percentages:\n", percentages)
```

```
Class distribution:
  LUNG_CANCER
YES      238
NO       38
Name: count, dtype: int64
Class percentages:
  LUNG_CANCER
YES    86.231884
NO    13.768116
Name: count, dtype: float64
```

- 238 Pos case of LC (86%)
- 38 Neg case of LC (14%)

- When we look at class balances we can see a strong class imbalance exist.

```
# Compare mean ages
mean_ages = data3.groupby('LUNG_CANCER')['AGE'].mean()
print("Mean ages by group:\n", mean_ages)

# T-test for age difference
from scipy.stats import ttest_ind
ages_yes = data3[data3['LUNG_CANCER'] == 1]['AGE'].dropna()
ages_no = data3[data3['LUNG_CANCER'] == 0]['AGE'].dropna()
t_stat, p_val = ttest_ind(ages_yes, ages_no, equal_var=False)
print(f"T-test: t={t_stat:.2f}, p={p_val:.4f}")

Mean ages by group:
  LUNG_CANCER
0    60.684211
1    63.264706
Name: AGE, dtype: float64
T-test: t=1.55, p=0.1285
```

- Patient without lung cancer have an avg age of 60.7
- Patient with lung cancer have an avg age of 63.26
- t-statistic = 1.55
- p-value = 0.1285
  - Age difference is not statistically sig.
  - Cannot reject the null hypothesis that the age distributions are the same
- This suggests that while lung cancer patients are slightly older on average in your dataset (63.3 vs 60.7 years), age alone is not a strong differentiating factor

```
[40]: categorical_columns2 = ["GENDER", "SMOKING", "YELLOW_FINGERS", "ANXIETY", "PEER_PRESSURE",
                             "CHRONIC_DISEASE", "FATIGUE", "ALLERGY", "WHEEZING", "ALCOHOL_CONSUMING",
                             "COUGHING", "SHORTNESS_OF_BREATH", "SWALLOWING_DIFFICULTY", "CHEST_PAIN"]

for col in categorical_columns2:
    data3[col] = data3[col].astype("category").cat.codes

corr_matrix = data3.corr(numeric_only=True)
#print(corr_matrix["LUNG_CANCER"].sort_values(ascending=False))

[48]: #corr_matrix.info()

[44]: print(corr_matrix["LUNG_CANCER"].sort_values(ascending=False))

LUNG_CANCER      1.000000
SYMPTOM_SCORE    0.397805
ALLERGY          0.333552
ALCOHOL_CONSUMING 0.294422
SWALLOWING_DIFFICULTY 0.268940
COUGHING         0.253027
WHEEZING         0.249054
PEER_PRESSURE    0.195086
CHEST_PAIN       0.194856
YELLOW_FINGERS   0.189192
FATIGUE          0.160078
ANXIETY          0.144322
CHRONIC_DISEASE  0.143692
AGE              0.106305
SHORTNESS_OF_BREATH 0.064407
GENDER           0.053666
SMOKING          0.034878
Name: LUNG_CANCER, dtype: float64
```

## • Correlation Analysis


- Allergy shows the 2nd highest correlation with lung cancer diagnosis (0.33)
- Smoking as a very weak correlation.
  - VERY SURPRISING!!!!
- SMOKING shows near-zero correlation (0.03) (NEED TO BE INVESTIGATED FURTHER)

```

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(12,10))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", square=True)
plt.title("Correlation Matrix for Lunch Cancer Dataset")
plt.show()

```



- Dark Red = Strong Pos
- Dark Blue = Strong Neg

## ANALYSIS RESULTS

- Class Distribution Analysis
  - Yes: 238 cases (86.2%)
  - No: 38 (13.8%)
  - Severe imbalance.
- Age Comparison Between Groups
  - Yes mean: 63.26 +/- 8.1 yrs
  - No mean: 60.68 +/- 9.8 yrs
- Welch's t-test
  - $t = 1.55$
  - $p = 0.1285$ 
    - Not Stat sig
- Symptom Score Relationship
  - Kendall's Tau
    - 0.34 and  $p = < 0.001$
  - Strong pos association between symptom score and cancer likelihood.
- Feature Correlation Analysis
  - Point=Biserial Correlation
    - Symptom\_score = 0.40
    - Allergy = 0.33
    - Alcohol consuming = 0.29
- Categorical Feature Analysis
  - Chi-Square Test

